

# Graphical Abstract

## **StarPlat: A Versatile DSL for Graph Analytics**

Nibedita Behera, Ashwina Kumar, Ebenezer Rajadurai T, Sai Nitish, Rajesh Pandian M, Rupesh Nasre

## Highlights

### **StarPlat: A Versatile DSL for Graph Analytics**

Nibedita Behera, Ashwina Kumar, Ebenezer Rajadurai T, Sai Nitish, Rajesh Pandian M, Rupesh Nasre

- StarPlat, a DSL for graph analytics which allows users to write a high level algorithmic specification of their static graph processing, which captures the parallelism intentions but decouples the target architecture.
- An intermediate representation common across backends that captures all the essential information associated with constructs for their translation in the target architecture.
- A code-generation scheme to support translation of the intermediate representation into efficient codes for the multicore (using OpenMP), distributed (using OpenMPI), and many-core (using CUDA) backends.
- Performance analysis of the generated BC, PR, SSSP, and TC codes for different target hardware on a variety of popular graphs, and illustrating competitive performance against that of the hand-crafted frameworks.

# StarPlat: A Versatile DSL for Graph Analytics

Nibedita Behera, Ashwina Kumar, Ebenezer Rajadurai T, Sai Nitish,  
Rajesh Pandian M, Rupesh Nasre

*<sup>a</sup>Department of Computer Science and Engineering, Indian Institute of Technology  
Madras, , Chennai, 600036, Tamil Nadu, India*

---

## Abstract

Graphs model several real-world phenomena. With the growth of unstructured and semi-structured data, parallelization of graph algorithms is inevitable. Unfortunately, due to inherent irregularity of computation, memory access, and communication, graph algorithms are traditionally challenging to parallelize. To tame this challenge, several libraries, frameworks, and domain-specific languages (DSLs) have been proposed to reduce the parallel programming burden of the users, who are often domain experts. However, existing frameworks to model graph algorithms typically target a single architecture. In this paper, we present a graph DSL, named StarPlat, that allows programmers to specify graph algorithms in a high-level format, but generates code for three different backends from the same algorithmic specification. In particular, the DSL compiler generates OpenMP for multi-core systems, MPI for distributed systems, and CUDA for many-core GPUs. Since these three are completely different parallel programming paradigms, binding them together under the same language is challenging. We share our experience with the language design. Central to our compiler is an intermediate representation which allows a common representation of the high-level program,

---

*Email addresses:* `cs20s023@cse.iitm.ac.in` (Nibedita Behera ),  
`cs20d016@cse.iitm.ac.in` (Ashwina Kumar), `ebenezerrajadurai5@gmail.com`  
(Ebenezer Rajadurai T), `bsainitishkumar@gmail.com` (Sai Nitish),  
`mrprajesh@cse.iitm.ac.in` (Rajesh Pandian M), `rupesh@cse.iitm.ac.in` (Rupesh Nasre)

*URL:* <https://orcid.org/0000-0002-1563-8686> (Nibedita Behera ),  
<https://orcid.org/0000-0001-6425-7479> (Ashwina Kumar),  
<https://orcid.org/0000-0003-4702-4678> (Rajesh Pandian M),  
<https://orcid.org/0000-0001-7490-625X> (Rupesh Nasre)

from which individual backend code generations begin. We demonstrate the expressiveness of StarPlat by specifying four graph algorithms: betweenness centrality computation, page rank computation, single-source shortest paths, and triangle counting. Using a suite of ten large graphs, we illustrate the effectiveness of our approach by comparing the performance of the generated codes with that obtained with hand-crafted library codes. We find that the generated code is competitive to library-based codes in many cases. More importantly, we show the feasibility to generate efficient codes for different target architectures from the same algorithmic specification of graph algorithms.

*Keywords:* Graph Algorithms, Domain-Specific Language, OpenMP, MPI, CUDA

---

## 1. Introduction

The graph data structure has become an integral component of many real-world applications for modelling relationships in their data today. Enormous growth of unstructured and semi-structured data has led to these graphs growing to billions of edges. Therefore, parallel graph analytic solutions are inevitable to scale to such large graph sizes. The last two decades have witnessed significant advances in hardware towards parallel processing, which has also resulted in major developments in the software support, be it in the form of libraries or programming languages. These hardware and software architectures are primarily suited for regular codes wherein the data access, control flow, and communication patterns are statically identifiable. As an example, tiling of regular matrix computations can now be performed automatically by the compiler for improved cache benefits or with minimal communication.

Unfortunately, existing widely-used compilers perform poorly in the case of graph algorithms. This is due to the inherent *irregularity* in sparse graph processing, wherein the access patterns are dependent on the input (which is unavailable at compile time). On the other hand, our community has shown that graph algorithms exhibit enough parallelism to keep the cores busy on several real-world graphs [? ]. Unfortunately, manually exploiting this parallelism on various hardware is quite challenging. The programmer needs to be an expert in the application domain to exploit algorithmic properties, in high-performance computing (HPC) to model the computation to suit

the target hardware, and also in computing systems to optimize the overall application and the support software on the given infrastructure. Even if corporations and institutions can find and afford such an expert, the expert is unlikely to be a best-fit for another application domain. A sore practical reality is that we have domain experts who are not HPC or systems experts, and we have HPC experts who may not know enough about the underlying application domain.

A viable alternative towards improved productivity is a graph library or a domain-specific language, which allows domain experts to express their algorithm using API or high-level constructs, and the library or the DSL compiler taking care of generating high-performing code for the target hardware. A range of such frameworks exists today [? ? ? ? ]. Multiple of these frameworks partially or fully hide the parallelism intricacies, provide mnemonics for scheduling strategies, and perform program analysis to identify races to generate synchronization code. It is often seen that the amount of code one needs to write reduces considerably compared to that in a hand-crafted explicitly parallel code.

One of the limitations of the existing frameworks (libraries or graph DSLs) is that they target a specific hardware architecture. For instance, Ligra [? ] is a graph processing frameworks specifically for shared memory systems. Greenmarl [? ] primarily targets multicore devices and supports distributed implementation through Pregel API. Gunrock [? ] is a CUDA library for graph processing on GPUs. GraphIt [? ] is a DSL that targets shared memory and manycore systems. Rarely, a framework targets more than one backends (e.g., Kokkos [? ]). For the best performance and considering the range of HPC hardware we use today, libraries often restrict themselves to a certain target. Multi-core parallelism permits data shared across all the running threads, while a cluster-level parallelism demands explicit communication, while further, many-core parallelism necessitates a hierarchical computation and SIMD execution for the best performance. While domain-experts are aware of these basic differences, they should be able to *specify* rather than *code up* these for an architecture to achieve desired performance. Therefore, it is ideal if the domain-specific language encompasses different backend peculiarities in its design. This bears the advantage of simplified and efficient code generation, and avoids *patching* a language originally designed for a different parallel architecture.

In this work, we propose a graph DSL named StarPlat which allows a user to provide an algorithm specification of graph problems using its high-




Figure 1: Overview of StarPlat compiler

level graph specific constructs and generates code for multiple backends from the same algorithmic specification (currently, multi-core, distributed, and many-core). The constructs are carefully designed to abstract the parallelization specific implementations from users, while ensuring generation of high performing graph codes. The StarPlat compiler moulds the high level information embedded in these constructs to different architecture specific implementations. The compiler translates the DSL code to C/C++. It uses OpenMP parallelism for the multicore setting, MPI for the distributed setup, and CUDA for the many-core architecture. Encompassing these three very different backends in the same DSL compiler is challenging, and we highlight these challenges in this manuscript. We proudly admit that we build upon the insights provided by the existing libraries and DSLs, borrow certain constructs from these frameworks (while providing our own), and generate code competitive to these frameworks in terms of performance. In particular, we illustrate the versatility of StarPlat with a discussion on four graph algorithms: Betweenness Centrality (BC), PageRank (PR), Single Source Shortest Paths (SSSP), and Triangle Counting (TC).

This work makes the following technical contributions:

- StarPlat<sup>1</sup>, a DSL for graph analytics which allows users to write a high

<sup>1</sup><https://github.com/nibeditabh/StarPlat>

81 level algorithmic specification of their static graph processing, which  
82 captures the parallelism intentions but decouples the target architec-  
83 ture.

- 84 • An intermediate representation common across backends that captures  
85 all the essential information associated with constructs for their trans-  
86 lation in the target architecture.
- 87 • A code-generation scheme to support translation of the intermediate  
88 representation into efficient codes for the multicore (using OpenMP),  
89 distributed (using OpenMPI), and many-core (using CUDA) backends.
- 90 • Performance analysis of the generated BC, PR, SSSP, and TC codes  
91 for different target hardware on a variety of popular graphs, and il-  
92 lustrating competitive performance against that of the hand-crafted  
93 frameworks.

94 The rest of the article is organized as follows: Section 2 presents the lan-  
95 guage specification and the intermediate representation. Section 3 describes  
96 the code-generation scheme followed for the translation of the DSL code for  
97 each backend. Section 4 provides an overview of the backend-specific opti-  
98 mizations StarPlat employs for efficient code generation. The experimental  
99 evaluation of the generated code for each backend is discussed in Section 5.  
100 Section 6 discusses the related work for graph analytics. We summarise our  
101 experience and conclude in Section 7.

## 102 2. StarPlat Language and Frontend

103 The high-level philosophy of StarPlat is to relieve the user of the paral-  
104 lelization constructs as much as possible, and to achieve performance compet-  
105 itive to hand-tuned codes by providing constructs and hints on aggregates.  
106 In rare cases, when it is a must to have a trade-off between abstraction and  
107 performance, we have taken a conscious decision to prioritize abstraction.  
108 This is because the language is meant primarily for domain-experts (rather  
109 than HPC experts).

110 From day one, the language was designed to abstract away the hardware.  
111 This was a challenge, since the backends are quite different. But we have  
112 found commonalities at the algorithmic level which we encode using specific  
113 constructs, which could be then translated to the appropriate backend code.

114 For instance, when a lock-based synchronization was required in OpenMP,  
 115 it also demanded communication in MPI, and a lock-free synchronization in  
 116 CUDA.

117 StarPlat provides various abstractions and data types relevant to static  
 118 graph algorithm, such as **Graph**, **node**, **edge**, **propNode** (for node property)  
 119 etc. The programmer writes in a procedural style and hints at the paralleliza-  
 120 tion opportunities using aggregate iteration constructs such as **forall** along  
 121 with other inherently parallel operations. The compiler decides whether to  
 122 exploit this parallelization (e.g., nested **forall**). Several algorithms can be  
 123 viewed as iterative procedures repeatedly executed until convergence. The  
 124 convergence criteria is application dependent. Such an iterative procedure  
 125 is described using a **fixedPoint** construct. StarPlat follows a "batteries  
 126 included" approach and has several utility functions for the data types. In-  
 127 ternally, each of the functions is implemented target-hardware-wise.

## 128 2.1. Compiler Overview

129 As in a general-purpose compiler, StarPlat's compiler is split into fron-  
 130 tend and backend, as shown in Figure 2. The frontend is responsible for  
 131 ensuring syntactic and semantic consistency of the StarPlat implementation.  
 132 The frontend builds an abstract syntax tree (AST) of the supplied input  
 133 code. The AST is common across all the backends and is populated with the  
 134 metadata for each construct during the parsing stage of the compiler. It is  
 135 then fed to the appropriate code generator depending upon the target code  
 136 the user wishes to generate.

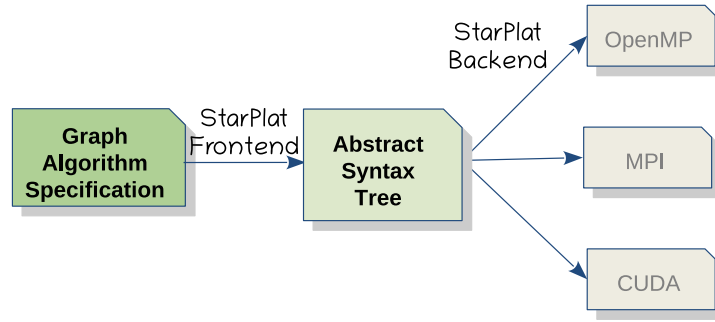


Figure 2: Process flow of the StarPlat compiler

## 137 2.2. Language Overview with an Example: SSSP

138 Single source shortest path computation (SSSP) finds a shortest path  
 139 from a designated source node to every other node in the graph. Parallel



SSSP computation can be modelled as an iterative procedure, as illustrated in Figure 3. This is a variant of the Bellman-Ford's algorithm.

```

142 1 function computeSSSP (Graph g, node src) {
143 2     propNode<int> dist;
144 3     propNode<bool> modified;
145 4     g.attachNodeProperty(dist = INF, modified =
146     False);
147 5     src.modified = True;
148 6     src.dist=0;
149 7     bool finished = False;
150 8     fixedPoint until (finished: !modified) {
151 9         forall (v in g.nodes().filter(modified
152         = True)) {
153 10             forall (nbr in g.neighbors(v)) {
154 11                 edge e = g.getEdge(v, nbr);
155 12                 <nbr.dist, nbr.modified> =
156                 <Min(nbr.dist, v.dist +
157                 e.weight), True>;
158 13     } } } }
```

Fig 3: SSSP computation in StarPlat

The function `compute_SSSP` takes two arguments: a graph `g` and source vertex `src`. Statement 2 declares a variable `dist` of type `propNode`. `propNode<int>` marks `dist` as an integer type attribute associated with each vertex. Statement 3 creates a `bool` type vertex attribute `modified`. In Statement 4, the function `attachNodeProperty` associates the two attributes with each vertex of `g`, and initializes the property values of `dist` to infinity (`INT_MAX`) and `modified` to false, as aggregate operations. In Statements 5–6, the `dist` and `modified` attributes for vertex `src` are assigned specific values. In particular, the distance of `src` is set to zero, and the vertex is also marked as modified. The `fixedPoint` construct in Statement 8 makes sure the iterative procedure executes until the variable `finished` becomes true. The variable `finished` is updated in each iteration, and decides if another iteration is necessary. A fixed-point is reached when no vertices are modified. This is achieved by or-ing the `modified` attribute of each node and assigning to `finished`. The iteration construct `forall` in Statement 9 specifies parallel iteration over the graph nodes. In addition, the optional `filter` clause allows processing a subset of the vertices meeting a criterion. The `forall` construct in Statement 10 specifies

176 iterating over all the neighbors of vertex `v` in parallel and for each neighbor  
 177 `nbr`, the corresponding edge is relaxed (Statements 11–12). The multiple  
 178 assignment construct in Statement 12 performs multiple tasks and is a suit-  
 179 able example of the abstraction provided by the DSL’s constructs (borrowed  
 180 from Green-Marl [? ]). The rvalue `Min(nbr.dist, v.dist + e.weight)` is  
 181 assigned to the `dist` attribute of `nbr` with the sum of `v.dist` and `e.weight`  
 182 based on whether the alternative distance via `v` is smaller than the existing  
 183 distance of `nbr`. If `nbr.dist` is updated, the `modified` attribute of `nbr` is set.  
 184 Such a construct allows processing related statements together (as a critical  
 185 section, to be precise) and translates to an update logic protected by some  
 186 form of synchronization to prevent data races.

187 We highlight that the SSSP code structurally resembles the algorithm  
 188 specified in a textbook, with a few changes. Hence, StarPlat eases imple-  
 189 mentation, analysis, and modification of graph algorithms from the context  
 190 of parallelization. The multi-core translation of this code would follow a  
 191 similar structure as the StarPlat code with OpenMP pragmas inserted; the  
 192 distributed version differs in terms of complexities involved in scattering and  
 193 gathering data across devices; while a many-core version need to pre-transfer  
 194 the data being used inside `forall` from CPU to GPU.

195 We now delve deeper into various StarPlat constructs.

### 196 2.3. Language Constructs

197 We first discuss various data types, followed by iteration schemes and  
 198 reductions.

#### 199 2.3.1. Data Types

200 StarPlat supports graph theoretic concepts as first-class types such as  
 201 Graph, node, edge, node attribute, edge attribute, etc. It also supports  
 202 primitive data types: `int`, `bool`, `long`, `float`, and `double`.

203 The **Graph** data type encapsulates the operations and properties of a  
 204 standalone graph. The properties include its nodes, edges, number of nodes,  
 205 number of edges, etc. StarPlat stores the graph in Compressed Sparse Row  
 206 (CSR) format, which provides the storage benefits of adjacency lists, and  
 207 also allows seamless transfer across devices, due to the use of offsets. The data  
 208 type also facilitates information gathering and manipulation at the node and  
 209 the edge levels. Since the nodes and edges are tightly bound to the graph,  
 210 it becomes convenient for Graph to support this through various library  
 211 functions. For instance, as per the semantics of `neighbors(u)`, it returns

the outgoing neighbors in case of a directed graph and all the neighbors in case of an undirected one. For directed graphs, graph type also exposes a function that returns the incoming neighbours to a node, `nodesTo()`. This needs Graph to maintain a CSR representation for also its transpose, which becomes handy in algorithms which perform computation on a transposed version of the input graph (e.g., Betweenness Centrality).

```

1  function foo(Graph g, propNode<int> modified ,
2      propNode<int> data) {
3      SetN<g> nodeSet;
4      for(v in g.nodes().filter(modified)) { //
5          Sequential loop
6          for(nbr in g.neighbors(v)) {
7              if(nbr.data > THRESHOLD) {
8                  nodeSet.addNode(nbr);
9              }
10         }
11     }
12 }
```

Fig 4: Example program to illustrate various data types in StarPlat

A node and an edge can in themselves have properties associated with them. In the SSSP problem setting, a vertex's distance can be viewed as a node property, being computed by the corresponding algorithm. Similarly, in the BC computation, the betweenness centrality values of each node can be viewed as a property. The `propNode` datatype facilitates declaring property for nodes of a graph with the provision of specifying its type. The `attachNodeProperty` function provided by the Graph type binds this property to the graph's nodes and initializes the property values if provided. Line 2 in our SSSP code from Figure 3 specifies the declaration of a node property `dist` of type `int`. The `attachNodeProperty` binds `dist` to the graph and optionally, initializes the distance attribute for each vertex (e.g., to infinity in Figure 3). Similarly, `propEdge` datatype is associated with edges, and has otherwise the same semantics as that of `propNode`. The `attachEdgeProperty` function binds the property to the graph's edges.

StarPlat also provides collection types such as `List`, `SetN`, and `SetE`. `List` allows the presence of duplicates whereas `SetN` and `SetE` store unique nodes and edges respectively. Line 2 in Figure 4 shows an example usage. The separation of sets between nodes and edges enables choosing the relevant implementation in vertex-based vs. edge-based codes.

### 246 2.3.2. Parallelization and Iteration Schemes

247 **forall** is an aggregate construct in StarPlat which can process a set of  
 248 elements in parallel. Its sequential counterpart is a simple **for** statement.  
 249 Currently, StarPlat supports vertex-based processing<sup>2</sup>. The parallel **forall**  
 250 supports various ranges it can iterate on (e.g., nodes in the whole graph or  
 251 neighbors of a node, as shown in Figure 3, Lines 9 and 10).

252 The function **g.nodes()** called on a graph **g** returns a sequence of nodes  
 253 which can be iterated upon. To iterate over the neighbors of a node **u**,  
 254 the functions **g.neighbors(u)**, **g.nodesTo(u)** and **g.nodesFrom(u)** return  
 255 a similar sequence. The **forall** body can be executed selectively for the  
 256 nodes satisfying a certain boolean expression based on the node label or,  
 257 node’s property by including a **filter** construct. Line 3 in Figure 4 shows  
 258 its usage using the node property **modified**.

259 Considering that several graph algorithms can be well represented using  
 260 a single outer parallel loop, and that the analysis of nested parallel loops gets  
 261 complicated, currently, StarPlat supports only outer level parallelism. Hence  
 262 a nested **forall** in the DSL results in a parallel outer loop and a sequential  
 263 inner loop in the target code. One may argue that an outer loop over vertices  
 264 and an inner loop over neighbors can benefit from nested parallelism, and we  
 265 agree. However, (i) such a processing can be well taken care of by an edge-  
 266 based parallelism (to be supported), and (ii) since we target large graphs,  
 267 even a vertex-based processing has enough parallelism to keep the resources  
 268 busy.

269 For a graph processing DSL, traversals become the fundamental building  
 270 blocks. StarPlat provides breadth-first traversal as a construct, borrowed  
 271 from GreenMarl [? ].

272 **iterateInBFS(v in g.nodes() from root) {...}**  
 273 **iterateInBFS** performs a BFS traversal of the graph from the given root  
 274 node. The underlying processing is level-by-level and iterates in parallel over  
 275 the visited nodes in a specific level. On visiting a node in a level, it executes  
 276 the body statements and forms the next set of visited nodes. The **filter()**  
 277 construct can be utilized to explore the neighborhood of a visited node selec-  
 278 tively. Similarly, **iterateInReverse** performs a reverse BFS traversal in a

---

<sup>2</sup>We are adding support for edge-based processing, which needs changes to the under-  
 lying data representation. Compressed Sparse Row (CSR) storage format is suited for  
 vertex-based processing.

level-synchronous manner and extracts parallelism at each level in the computation of the body statements. Note that `iterateInBFS` is a prerequisite to use `iterateInReverse`, since the former builds the BFS DAG to be traversed through in the latter. The functions `neighbors`, `nodesTo` and `nodesFrom` have a subtle change in their meaning when used inside `iterateIn...` constructs: they correspond to the neighbors in the BFS DAG rather than the original graph `G`. We have found this semantics change to satisfy our natural inclination to write a code. We illustrate it in the BC computation (Appendix Appendix A).

### 2.3.3. Reductions

Reductions are one of the popular parallel programming primitives, and can be useful in achieving efficient computation. Specifying a reduction in the DSL also helps in conveying a necessity of synchronization. Unfortunately, it does not directly fit into the philosophy of StarPlat design to support reduction as a language construct. Therefore, as a golden-mid, StarPlat permits usage of certain relative C-operators (e.g., `+=`) to convey reduction. This "trick" allows us to retain the abstraction and still achieve efficiency of the generated code. The reduction operators supported by StarPlat are tabulated in Table 1.

Operator	Reduction Type
<code>+=</code>	Sum
<code>*=</code>	Product
<code>++</code>	Count
<code>&amp;&amp;=</code>	All
<code>  =</code>	Any

Table 1: Reduction operators in StarPlat

We illustrate the usage of reduction in Figure 5. The introduction of reduction (Line 7 in the code makes sure the `accum` variable has a deterministic result at the end of the parallel region. Note that Line 5 involves a thread-local variable `count` and does not need reduction. On the other hand, if nested parallelism was supported, `count` would also need a reduction. The reduction operators in StarPlat translate to library based implementations

```

304 of reduction in the target backend.3
305     1 int accum = 0;
306     2 forall(v in g.nodes()) {
307     3     int count = 0;
308     4     forall(nbr in g.neighbors(v)) {
309     5         count = count + nbr.A;    // regular
310                                     assignment
311     6     }
312     7     accum += count; //sum reduction in assignment
313                             form
314     8 }

```

Fig 5: Reduction example

#### 315 2.3.4. *fixedPoint and Min/Max Constructs*

316 Several solutions to graph algorithms are iterative, and converge based  
317 on conditions on node attributes. StarPlat provides a `fixedPoint` construct to  
318 specify this succinctly. Its syntax involves a boolean variable and a boolean  
319 expression on node-properties forming the convergence condition, as shown  
320 below.

```

321 fixedPoint until (var: convergence expression) {...}

```

322 Line 8 of Figure 3 in SSSP’s specification uses the `fixedPoint` construct to  
323 define the convergence condition. The loop iterates till at least one node’s  
324 `modified` property is set to true.

325 StarPlat provides constructs `Min` and `Max` which perform multiple assign-  
326 ments based on a comparison criterion. This can be useful in update-based  
327 algorithms like SSSP, where an update on node properties is carried out on  
328 a desired condition, while taking care of potential data races.

329 In the SSSP computation, the `Min` construct is used to encode the relax-  
330 ation criteria in Line 12. The neighboring node’s distance is updated if the  
331 alternative distance via vertex `v` is smaller than `nbr`’s current distance. The  
332 update of the `dist` property based on this comparison specified using `Min`  
333 results in an update of the `modified` property to `True`.

334 StarPlat also has aggregate functions `minWt` and `maxWt` to find the mini-  
335 mum and maximum edge weights.

---

<sup>3</sup>Currently, the onus of writing the correct operator is onto the DSL user. We are adding program analysis to StarPlat which would relieve the user of worrying about this.

## 336 2.4. Abstract Syntax Tree (AST)

337 Each meaningful non-terminal that builds the language are at the highest  
338 level denoted as **ASTNode**. The **ASTNode** is the parent class of all the nodes  
339 that are part of the abstract syntax tree. The **ASTNode** forges to child class  
340 nodes that signify specialization like **statement** Node, **Expression** Node.  
341 The **statement** node is the parent class of all statement types like assign-  
342 ment, declaration, control flow, procedure call, etc. Each of these nodes can  
343 sometimes be composed of other nodes, owing to the way the construct is  
344 defined in the parser. For instance, **forallStmt**'s construct is composed of  
345 a body, an iteration space which is defined by an iterator and range (func-  
346 tions) and an optional filter expression on the iterated items. The body in  
347 **forallStmt** is represented as a statement node, iterator as an **Identifier**  
348 node, ranges as **proc\_callExpr**, and filter expressions as *Expression* node in  
349 the AST. The other node types are **Identifier**, **iterateReverseBFS**, **Type**,  
350 **reductionCall**, **fixedPointStmt**, **formalParam** etc.

351 Most of the node's data are populated in the parsing stage itself. Data re-  
352 lated to the type of the symbols are added during an additional pass through  
353 the already built AST. Given the target backend, the construct's AST are  
354 populated with additional information relevant for the code-generation in  
355 that specific backend. The analyzer phase before the code generation per-  
356 forms analysis for optimized and efficient code generation. Some of these  
357 optimizations are specific to the target backend. The separation of nodes  
358 into different subclasses allows maintaining handlers for different node types  
359 in the code generator. The code-generator forms the backend part of the com-  
360 piler. Unlike the frontend which was primarily target-independent, backend  
361 is tightly coupled with the target parallelization platform.

## 362 3. StarPlat Code Generator

363 The code generation phase supports the generation of efficient code for  
364 primarily three target backends, shared memory, distributed, and many-core  
365 platforms. The generated code for each backend also makes use of specific  
366 backend libraries to handle the parallel processing intuitively specified in  
367 the StarPlat code. Shared memory implementation uses OpenMP as the  
368 threading library, distributed memory implementation uses MPI, and many-  
369 core implementation uses CUDA.

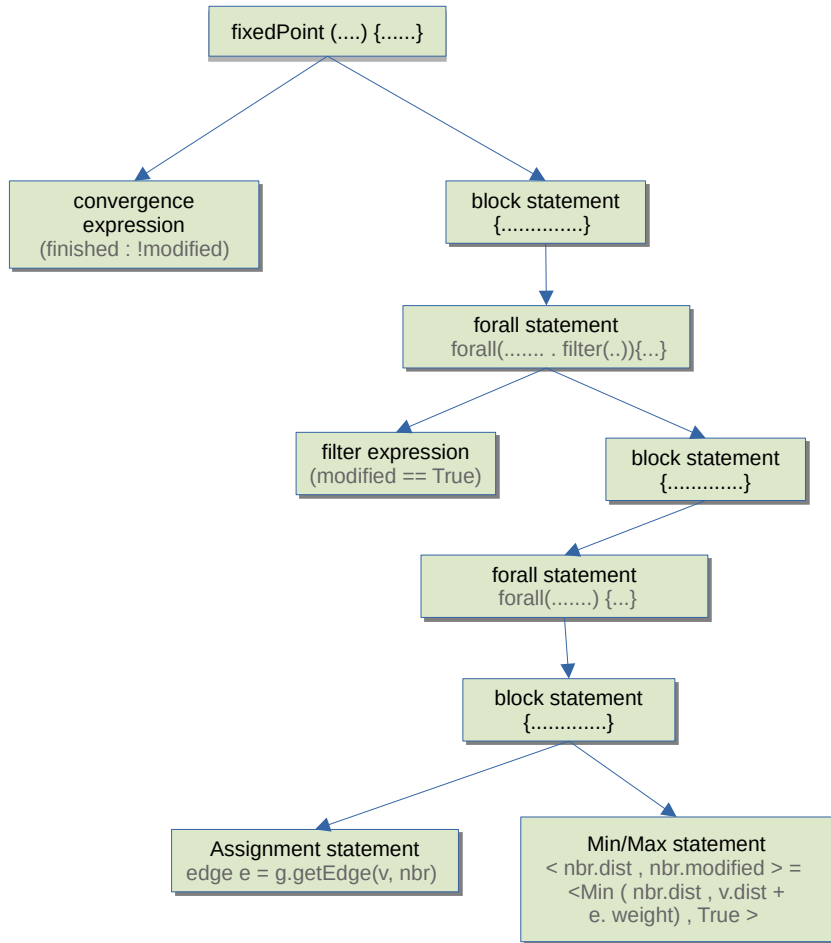


Figure 6: AST for SSSP's `fixedPoint` construct in Figure 3



### 370 3.1. Graph Representation and Storage

371 While there are several formats in which graphs can be stored, we pre-  
372 ferred a format which:

- 373 • can preferably work across all the backends
- 374 • works well with vertex-centric algorithms, common in graph processing
- 375 • can be easily split

376 The last preference is not only due to the MPI backend, but also due to  
377 our future plans to extend StarPlat for multi-GPU and heterogeneous back-  
378 ends. Since offset-based formats fit our requirements, we had the options of  
379 compressed sparse row (CSR) and coordinate list (COO) formats, of which  
380 CSR fits the second preference. Therefore, we chose to use it. CSR is also  
381 popularly used in sparse graph processing frameworks. StarPlat has a back-  
382 end graph library which takes care of loading a graph from the input file,  
383 and storing it in the CSR format.

384 The OpenMP backend does not pose any issues with respect to storing  
385 the graphs. In case of the MPI backend, the graph is loaded by rank 0  
386 process and distributed using a simple block partitioning approach among  
387 the MPI processes, resulting in *local* and *remote* vertices for each process.  
388 The information of only the local vertices is available for direct access, while  
389 that of the remote vertices needs to be communicated with explicit send-  
390 receives. The MPI backend maintains a mapping between the local vertex id  
391 and its corresponding global vertex id, for correct processing and for sending  
392 and receiving data. This is because, for instance, a large global vertex id may  
393 still map to a local id of 0 in an MPI process. The CUDA backend needs  
394 to transfer the graph and the associated attributes fully from the host to  
395 the GPU (using multiple `cudaMemcpy` calls). Since the algorithms work with  
396 static graphs, only the modified attribute information needs to be brought  
397 back from the GPU to the host, at the end of the processing.

398 Seamless management of both the graph representation and the attribute  
399 storage across various backends relieves the domain-user of complex and er-  
400 roneous graph handling.

### 401 3.2. Overall Flow

402 In the OpenMP backend, the outermost `forall` gets translated to have  
403 `#pragma omp parallel for` as a header. Input graphs are parsed and stored

404 as C++ classes in the generated code, while the nodes and the edges are rep-  
 405 resented as integers (for their ids). The node properties and edge properties  
 406 are translated as arrays of types specified in the property declaration. We  
 407 have found that generating OpenMP code is relatively easier compared to the  
 408 other two backends due to (i) a single device, (ii) shared memory paradigm  
 409 wherein all the variables are present in the same memory address space for  
 410 all the threads, and (iii) pragma based processing which does not need to  
 411 change the sequential code.

412 The MPI backend generates code that follows a bulk-synchronous pro-  
 413 cessing (BSP) style, which involves a computation step followed by a com-  
 414 munication step in every iteration. During the computation step, various  
 415 MPI processes execute the main-loop code, update the attributes of their lo-  
 416 cal vertices, and mark those for sending in a send-buffer for the corresponding  
 417 global vertices. During the communication phase, processes receive these up-  
 418 dates from other processes, convert those to their local vertices, and update  
 419 the attributes. Similar to OpenMP, the node edge attributes / properties are  
 420 translated as arrays.

421 Unlike OpenMP and MPI, the CUDA backend needs to convert the  
 422 `forall` loop into a GPU kernel, which has a different scope. This makes cer-  
 423 tain variables unavailable, which need to be explicitly allocated (`cudaMalloc`)  
 424 and transferred (`cudaMemcpy`). In addition, the `forall` loop body may con-  
 425 tain local variables, which get translated to thread-local variables in CUDA.  
 426 To distinguish between these two kinds of variables, the CUDA backend  
 427 performs a rudimentary analysis of the AST.<sup>4</sup> Apart from each outermost  
 428 `forall`, the initialization to attribute values needs to be converted to a sep-  
 429 arate kernel (which is a loop in OpenMP and MPI). The kernel launch config-  
 430 uration is set to a fixed number of threads per block (1024 in the generated  
 431 code, which performs the best on an average in our experiments) and the  
 432 number of blocks proportional to the number of vertices.

433 Apart from this, the boilerplate code consists of calls to timing functions  
 434 to time the initialization, the fixed-point processing, and the data transfer  
 435 where applicable.

---

<sup>4</sup>A full-fledged program analysis of the algorithm at the AST level is an ongoing work,  
 which involves identifying races and addressing those with proper synchronization or com-  
 munication.

### 436 3.3. Neighborhood Iteration

437 Iterating over the neighborhood is typically nested inside iterating over  
 438 the graph vertices (possibly, with a filter, as in Figure 3). In the case of the  
 439 OpenMP backend, the generated code snippet is as below.

```

440 1 #pragma omp parallel for
441 2 for (int v = 0; v < g.num_nodes(); v++) {
442 3     if (!modified[v]) continue;
443 4     for (int edge = g.indexofNodes[v]; edge <
444         g.indexofNodes[v + 1]; edge++) {
445 5         int nbr = g.edgeList[edge];
446 6         ...
447 7     } }
```

Fig 7: OpenMP code generated for neighborhood iteration

448 The MPI code is very similar, but goes over a predefined range based on  
 449 the process rank.

```

450 1 mpi::communicator world;
451 2 myrank = world.rank();
452 3 np = world.size();
453 4
454 5 part_size = ceil(g.num_nodes() / (float)np);
455 6 startv = myrank * part_size;
456 7 endv = startv + part_size - 1;
457 8 ...
458 9 for (int v = startv; v <= endv; v++) {
459 10    for (int edge0 = local_index[v-startv];
460        edge0 < local_index[v-startv+1]; edge0++)
461        {
462 11        int nbr = local_edgeList[edge0];
463 12        ...
464 13    } }
```

Fig 8: MPI code generated for neighborhood iteration

465 The CUDA code needs to separate the kernel call and kernel processing.  
 466 The kernel call sets up the launch configuration (number of thread-blocks  
 467 and the size of each thread-block) and passes appropriate parameters (graph,  
 468 attributes, algorithm specific variables such as the source node in SSSP, and

other miscellaneous variables used internally). The kernel uses 1D thread id, and utilizes the CSR copied to the GPU to go over the neighbors. Note that unlike OpenMP and MPI, there is no loop over the vertices, which is handled by parallel threads with which the kernel is launched.

```

473     1  __global__ void computeSSSP (...) {
474     2      unsigned id = blockIdx.x * blockDim.x +
475           threadIdx.x;
476     3      ...
477     4      for (int edge = gpu.OA[id]; edge < gpu.OA[id
478           + 1]; edge++) {
479     5          int nbr = g.gpu_edgeList[edge];
480     6          ...
481     7      }
482     8      ...
483     9  }
484    10  ...
485    11  computeSSSP<<<nblocks, blocksize>>>(...);
486    12  cudaDeviceSynchronize();

```

Fig 9: CUDA code generated for neighborhood iteration

The loop bound functions change appropriately for directed vs. undirected graphs, and for in-neighbors (`g.nodesTo(v)`) which demand using reverse adjacencies (`g.revIndexofNodes[v]`).

### 3.4. Reductions

Recall the reduction supported in StarPlat (e.g., Line 7). In the OpenMP backend, the reduction on a scalar translates to OpenMP's `reduction` clause as below. OpenMP reduction implementation creates and updates a private copy of the reduction variable per thread. At the exit of the parallel region, the local changes are accumulated in the global variable. This significantly reduces the data race handling costs required for the correctness of operation on a shared variable.

```

498  #pragma omp parallel for reduction(+:accum)

```

The situation in CUDA gets complicated due to limitations on the number of resident thread-blocks, which can block the kernel. One option to counter this is to use a reduction from the host as a separate kernel (e.g., using `thrust::reduce`). But this demands terminating the kernel, coming back to the host, calling `reduce`, coming back to the host, and then calling

504 another kernel to perform the rest of the processing in the `forall` loop. This  
 505 not only adds complication to the code generation, but also makes the overall  
 506 processing inefficient. Therefore, we rely on atomics to generate the func-  
 507 tionally equivalent code (e.g., using `atomicAdd` in the example in Figure 5).  
 508 `atomicAdd(&accum, prop[nbr]);`

### 509 3.5. BFS Traversal

510 The `iterateInBFS` construct gets expanded to a parallel BFS. This single  
 511 construct expands to around 40 lines of OpenMP code, 60 lines of MPI code,  
 512 and 40 lines of CUDA code. Therefore, we show and discuss only crucial  
 513 parts of the snippets below.

514 The OpenMP backend maintains a vector of vectors for tracking vertices  
 515 level-wise as shown in the code snippet in Figure 10. The BFS has an outer  
 516 `while` loop and two inner loops: one to go over the vertices in a level (this is  
 517 parallelized at Line 6) and the other to add the explored vertices in a level to  
 518 the appropriate level-vector (Line 18), which acts as a frontier. The parallel  
 519 loop examines a vertex’s neighbors and checks if their current distances are  
 520 not set (value -1) and sets those to 1 + current vertex’s distance. To avoid  
 521 data races, this update needs to be done atomically using a CAS (Line 9).

```

522     1 levelNodes [ phase ]. push_back ( root );
523     2 ...
524     3 while ( bfsCount > 0 ) {
525     4     ...
526     5     #pragma omp parallel for
527     6     for ( int l = 0; l < prev_count ; l++ ) {
528     7         int v = levelNodes [ phase ] [ l ];
529     8         ...
530     9         dnbr =
531             __sync_val_compare_and_swap ( &bfsDist [ nbr ],
532             -1, bfsDist [ v ] + 1 );
533    10         if ( dnbr < 0 ) {
534    11             int num_thread = omp_get_thread_num ();
535    12             levelNodes_later [ num_thread ]. push_back ( nbr );
536    13         }
537    14         // body of iterateInBFS
538    15     }
539    16     phase = phase + 1;

```

```

540 17
541 18     for (int i = 0; i < omp_get_max_threads();
542        i++) {
543 19
544        levelNodes[phase].concat(levelNodes_later[i]);
545 20        ...
546 21    }
547 22 }

```

Fig 10: OpenMP code generated for the `iterateInBFS` construct

548 The MPI backend runs a similar outer `while` loop and goes over the  
549 vertices level by level, but only for the set of local vertices. Once a vertex's  
550 neighbor is updated, it is pushed to the send buffer (Line 14). At the end  
551 of each level, an all-to-all call exchanges the level values of local and remote  
552 vertices (Line 17). Different levels are separated by an MPI barrier in a BSP  
553 style (Line 23).

```

554 1 while(is_finished(startv, endv, active)) {
555 2     ...
556 3     while (active.size() > 0) {
557 4         int v = active.back();
558 5         active.pop_back();
559 6         ...
560 7         for (int j = local_index[v - startv]; j
561            < local_index[v - startv + 1]; j++) {
562 8             int w = local_edgeList[j];
563 9             //If local
564 10            if (d[w-startv] < 0) {
565 11                active_next.push_back(w);
566 12                d[w-startv] = d[v-startv] + 1;
567 13            } else
568 14                // send d[v], w, and v
569 15            ...
570 16        }
571 17        all_to_all(world, sendbuf, receivebuf);
572 18        for (int t = 0; t < np; t++) {
573 19            // process receives
574 20            if (d[w-startv] == d_v + 1)
575 21                //Body of iterateInBFS

```

```

576      22      }
577      23      MPI_Barrier(MPLCOMM_WORLD);
578      24      phase = phase + 1 ;
579      25  }

```

Fig 11: MPI code generated for the `iterateInBFS` construct

580 The CUDA backend poses a challenge for the `iterateInBFS` construct  
581 because unlike the other two backends, the code needs to be generated to  
582 be running on both the host and the device. The outer `do-while` loop runs  
583 on the host (Line 1), which internally calls the level-wise BFS kernel on the  
584 GPU (Line 13). Since the loop is on the host, for ensuring progress with the  
585 loop's condition, the flag (`finished`) needs to be repeatedly copied across  
586 devices (Lines 2 and 8), passed as a parameter to the kernel, and updated in  
587 the kernel whenever there is change of a vertex's level.

```

588      1  do {
589      2      H2D(...);
590      3
591      4      BFS<<<...>>>(...)
592      5      cudaDeviceSynchronize();
593      6      ++hops_from_source;
594      7
595      8      D2H(...);
596      9      ...
597     10
598     11 } while (!finished);
599     12
600     13 __global__ void BFS(...) {
601     14     ...
602     15     if (d_level[u] == *d_hops_from_source) {
603     16         ...
604     17         for (unsigned i = d_offset[u]; i < end;
605     18             ++i) {
606     19             unsigned v = ... if (d_level[v] ==
607     20                 -1) {
608     21                 d_level[v] = *d_hops_from_source
609     22                     + 1;
610     23                 *d_finished = false;
611     24             }

```

```

612      22          // Body of iterateInBFS
613      23 }      }      }

```

Fig 12: CUDA equivalent code for the `iterateBFS` construct

### 614 3.6. *Min/Max Constructs*

615 Recall the `Min` construct from SSSP (Line 12). In OpenMP, the condi-  
616 tional update on the node property is achieved through atomic implemen-  
617 tations of these operations built upon atomic CAS (Compare and Swap),  
618 made available as a StarPlat library. The generated OpenMP code for the  
619 `Min` construct in SSSP is as follows.

```

620      1  int dist_new = dist[v] + weight[e];
621      2  bool modified_new = true;
622      3
623      4  if (dist[nbr] > dist_new) {
624      5      int oldValue = dist[nbr];
625      6      atomicMin(&dist[nbr], dist_new);
626      7
627      8      if (oldValue > dist[nbr]) {
628      9          modified_nxt[nbr] = modified_new;
629     10 }      }

```

Fig 13: OpenMP code generated for the `Min` construct

630 In the MPI backend, the neighbor could be local or remote. If it is  
631 local, each process can update its property without synchronization. If it is  
632 remote, the property value to be updated is stored in the send buffer, which  
633 is eventually sent to the owning process (during the communication phase).

```

634      1  if (nbr >= startv && nbr <= endv) { // local
635      2      int e = edge0 + startv;
636      3      int dist_new = dist[v-startv] +
637      4      weight[e-startv];
638      5      bool modified_new = true;
639      6
640      7      if (dist[nbr-startv] > dist_new) {
641      8          dist[nbr-startv] = dist_new;
642      9          modified[nbr-startv] = modified_new;
643     10 }

```



```

644 10 } else { // remote
645 11     sendbuf[owner(nbr)][nbr] = dist[v-startv] +
646     weight[e-startv];
647 12 }
648 13 ...
649 14 all_to_all(world, sendbuf, receivebuf);

```

Fig 14: MPI code generated for the `Min` construct (comments are not generated)

The CUDA backend handles `Min/Max` constructs exactly similar to the OpenMP backend, except (i) the variables used are the GPU copies, and (ii) the atomic instructions are readily supported in CUDA (`atomicMin` and `atomicMax`).

```

654 1 int nbr = gpu_edgeList[edge];
655 2 int e = edge;
656 3 int dist_new = gpu_dist[v] + gpu_weight[e];
657 4
658 5 if (gpu_dist[nbr] > dist_new) {
659 6     atomicMin(&gpu_dist[nbr], dist_new);
660 7     gpu_modified_next[nbr] = true;
661 8     gpu_finished[0] = false;
662 9 }

```

Fig 15: CUDA code generated for the `Min` construct

### 3.7. *fixedPoint* Construct

The `fixedPoint` construct translates to a while loop conditioned on the fixed-point variable provided in the construct. Typically, the convergence is based on a node property, but it can be an arbitrary computation. This code is generated on the host, so it is similar in template for all the three backends. However, as before, the MPI code needs to gather this information for setting the flag `finished` from all the processes based on the vertex ids they operate on. Similarly, the CUDA code updates a copy of the `finished` flag on the GPU, which is `cudaMemcpy`'ed to the host.

```

672 1 while (!finished) {
673 2     finished = true;
674 3
675 4     if (...) { // Min-construct expansion

```

```

676      5      ...
677      6      finished = false; // fixedPoint
678                identifier updation
679      7  }    }

```

Fig 16: Code generated for the `fixedPoint` construct across all the backends

## 680 4. Optimizations

681 There are certain optimizations which a hand-crafted code can embed,  
682 which StarPlat’s code generator cannot do. For instance, if variants of an  
683 algorithm perform better on different backends, we cannot auto-generate the  
684 variants from the same algorithmic specification (e.g., push versus pull based  
685 processing [? ]), although the DSL is capable of specifying the variants  
686 separately. At the implementation level, persistent kernels in CUDA [? ]  
687 can improve execution time of algorithms on large diameter graphs. However,  
688 StarPlat generates the code for the `fixedPoint` construct on the host – across  
689 backends. Despite this, our observation has been that many optimizations  
690 can be readily embedded into the code generated by StarPlat:

- 691 • The algorithmic optimizations can be embedded into the DSL specifi-  
692 cation itself, which forms a common input to all the backends. As an  
693 illustration, Figure 21 in the Appendix presents a pull-based SSSP.
- 694 • A backend-specific optimization can be embedded into that specific  
695 backend, since such an optimization will not be explicitly mentioned in  
696 the DSL specification and will be abstracted away.

697 We discuss below the backend-specific optimizations.

### 698 4.1. OpenMP

699 **Avoiding false sharing in `iterateInBFS`.** In the case of BFS abstractions,  
700 nodes discovered by multiple threads at a particular level are pushed to a  
701 data structure required for graph traversal. The false sharing due to shared  
702 usage of the same array in this scenario has been alleviated by assigning local  
703 update space to each thread.

704 **Efficient fixed-point computation.** The `fixedPoint` construct converges  
705 on a specific condition on a single boolean node property. The change of  
706 convergence is tracked through a boolean fixed-point variable ideally needs

707 to be updated after analyzing the property values for all nodes. The update  
708 procedure has been optimized by updating the fixed-point variable along with  
709 the update to the property value for any node. Since, updates to a boolean  
710 variable by multiple threads are atomic by hardware, this does not lead to a  
711 performance loss.

712 **Using built-in atomics.** Atomics has proved to be lightweight. Hence,  
713 atomics implemented using built-in functions provided by the GCC compiler  
714 is being used in generated code for achieving exclusive access to a single  
715 statement by a thread instead of locks.

## 716 4.2. MPI

717 **Communication aggregation.** The communication aggregation opti-  
718 mization is used by the MPI code while updating the distance of remote  
719 neighbor vertices. Instead of sending multiple messages to a remote vertex  
720 by a processing node, a single message with local minimum value is sent to  
721 the remote vertex.

722 **Using Boost.** The packing/unpacking mechanism provided by Boost MPI  
723 library gives better performance in sending and receiving STL containers  
724 such as vectors, maps etc., and user-defined data types than MPI library  
725 calls.

726 **Quick index-based partitioning.** StarPlat currently uses an index based  
727 partitioning to distribute the graph among various MPI processes. Compared  
728 to using a partitioner such as Metis, ours has the down-side of increased com-  
729 munication due to several inter-partition edges. However, such a partitioning  
730 is quick, and also allows us to partition the graph arbitrarily based on the  
731 number of MPI processes. Presently, the implementation assumes that all the  
732 processes have equal number of vertices. Hence, we pad temporary vertices  
733 for the last process.<sup>5</sup>

## 734 4.3. CUDA

735 **Optimized host-device transfer.** We perform a rudimentary program  
736 analysis of the AST to identify variables that need to be transferred across  
737 devices. For instance, since graph is static, it need not be copied back from

---

<sup>5</sup>We definitely plan to exploit Metis and its parallel variants for MPI, multi-GPU, and heterogeneous backends.

GPU to CPU at the end of the kernel. In contrast, the modified properties need to be transferred back. Similarly, the `finished` flag is set on CPU, conditionally set on GPU, and read on the CPU again in the fixed-point processing. Therefore, the variable needs to be transferred to-and-fro. The `forall`-local variables are generated as device-only variables.

**Memory optimization in OR-reduction.** The way we write the `fixedPoint` construct, a `modified` property is used in computing the fixed-point. At a high-level, another iteration is necessary if any of the vertices’ `modified` flag is set. This is essentially a logical-OR operation. StarPlat takes advantage of this to generate a single flag variable which is set by threads in parallel (with dependence on hardware atomicity for primitive types). Managing this flag is cheaper than transferring arrays of the `modified` flags across devices, both in terms of time and memory.

## 5. Experimental Evaluation

To quantitatively assess StarPlat, we generate four popular algorithm implementations: Betweenness Centrality (BC), PageRank (PR), Single Source Shortest Path (SSSP), and Triangle Counting (TC). These are also coded and optimized using the other frameworks we compare against. The StarPlat code for SSSP is presented in Figure 3, while the other three are presented in the Appendix. Each code fits in about 30 lines, and with that effort, a domain-expert or a student can generate efficient implementations of the four algorithms for a multi-core setup, a distributed system, and a GPU.<sup>6</sup>

Since the DSL codes are rather short, the compilation is immediate. Therefore, we focus on analyzing the efficiency of the generated codes for the various backends. We compare the performance of the backend-specific code in StarPlat against the existing state-of-the-art graph analytic solutions (qualitative comparison is discussed in Section 6). StarPlat’s OpenMP backend is compared against Galois [? ], Green-Marl [? ], and Ligra [? ]; its MPI backend is compared against Gluon [? ]; and its CUDA backend is compared against GunRock [? ] and LonestarGPU [? ]. Galois is a C++-based framework for graph analytics, while Ligra is C-based. In addition, Galois also supports graph mutation, which is not the focus of other frameworks. Similar to StarPlat, Green-Marl is a graph DSL designed for the multi-core

---

<sup>6</sup>We believe this is remarkable, and would be appreciated by the community.

771 backend. Gluon is a distributed version of Galois built as a C++ framework.  
772 LonestarGPU is a manually optimized collection of CUDA codes, while Gun-  
773 rock is a CUDA-based library providing data-centric API for graph analytics,  
774 such as *advance*, *compute*, and *filter*. So, except for Green-Marl, we compare  
775 StarPlat-generated codes against manually optimized ones.

776 We use ten large graphs in our experiments, which are a mix of different  
777 types. Six of these are social networks exhibiting the small-world property,  
778 two are road networks having large diameters and small vertex degrees, while  
779 two are synthetically generated. One synthetic graph has a uniform random  
780 distribution (generated using Green-Marl’s graph generator), while the other  
781 one has a skewed degree distribution following the recursive-matrix format  
782 (generated using SNAP’s RMAT generator with parameters  $a = 0.57$ ,  $b =$   
783  $0.19$ ,  $c = 0.19$ ,  $d = 0.05$ ). They are listed in Table 2, sorted on the number  
784 of edges in each category. For unweighted graphs, we assign edge-weights  
785 selected uniformly at random in the range  $[1,100]$  (for SSSP).

Graph	Acronym	Num. Vertices (million)	Num. Edges (million)	Avg. Degree	Max. Degree
twitter-2010	TW	21.2	265.0	12	302,7
soc-sinaweibo	SW	58.6	261.0	4	4,0
orkut	OK	3.0	234.3	76.2813	33,3
wikipedia-ru	WK	3.3	93.3	55.4067	283,9
livejournal	LJ	4.8	69.0	28.257	22,8
soc-pokec	PK	1.6	30.6	37.5092	20,5
usaroad	US	24.0	28.9	2	
germany-osm	GR	11.5	12.4	2	
rm876	RM	16.7	87.6	5	128,5
uniform-random	UR	10.0	80.0	8	

Table 2: Input graphs

786 All our experiments, including the baselines we compare against, were  
787 run on IIT Madras AQUA cluster. The configuration of each compute node  
788 as follows: Intel Xeon Gold 6248 CPU with 40 hardware threads spread over  
789 two sockets, 2.50 GHz clock, and 192 GB memory running RHEL 7.6 OS.  
790 All the codes in C++ are compiled with GCC 9.2, with optimization flag  
791 -O3. Various backends have the following versions: OpenMP version 4.5,  
792 OpenMPI version 3.1.6, CUDA version 10.1.243 and run on Nvidia Tesla

793 V100-PCIE GPU with 5120 CUDA cores spread uniformly across 80 SMs  
 794 clocked at 1.38 GHz with 32 GB global memory and 48 KB shared memory  
 795 per thread-block.

### 796 5.1. *OpenMP*

797 Table 3 presents the running times of the codes in various frameworks  
 798 for the four algorithms. The execution time refers to only the algorithmic  
 799 processing and excludes the graph reading time, and is an average over three  
 800 runs (for each backend). The Galois framework failed to load the largest  
 801 graph TW (exited with a segfault). The framework also failed for BC com-  
 802 putation on 150 sources for US graph and exited by giving a PTS out of  
 803 memory error.

804 **Betweenness Centrality.** BC resembles all-pairs shortest paths, which  
 805 can be implemented by running SSSP from each vertex as a source. Com-  
 806 plete execution of BC on large graphs takes several hours, sometimes days.  
 807 Depending upon how an implementation stores BC results for each source,  
 808 the memory requirement can also increase proportional to  $\sim V \sim$  per source.  
 809 Therefore, literature presents results of running BC from one or only a few  
 810 vertices. We tabulate results for BC for number of source vertices as  $\{20,$   
 811  $80, 150\}$ . The source list for each graph is generated using a random gener-  
 812 ator and fixed across frameworks for consistent comparison. StarPlat’s BC  
 813 code outperformed Galois code for the road network graphs and some of the  
 814 social network graphs. Galois framework involves a scheduling policy that  
 815 ensures better load balance among threads. But for jobs with even work-  
 816 loads, the scheduling adds to the runtime overhead. Ligra has an optimized  
 817 forward BFS pass that bypasses distance computation for nodes, and com-  
 818 putes only the number of shortest paths for each node, thus saving on the  
 819 synchronization required for concurrent distance computation. This opti-  
 820 mization results in Ligra outperforming StarPlat and Galois. Interestingly,  
 821 the Green-Marl-generated code outperforms the other three frameworks on  
 822 most inputs. Green-Marl’s forward BFS traversal is optimized and chooses  
 823 the nature of traversal at each level based on specific criteria.

824 **PageRank.** StarPlat-generated PR code is competitive to that of Green-  
 825 Marl across inputs. Ligra, interestingly, has considerably slower implemen-  
 826 tation and takes significantly longer. This is primarily due to the loop sep-  
 827 aration between the computation of the PR values and the difference of  
 828 successive PR values for each vertex. Overall, Galois (about  $2\times$  faster than

Algo.	Framework	TW	SW	OK	WK	LJ	PK	US	GR	RM	UR	
BC	20	Galois	-	<b>1.840</b>	8.272	3.377	3.209	1.477	38.488	15.368	<b>2.861</b>	29.52
		Ligra	11.200	13.130	5.936	3.470	4.060	<b>1.220</b>	<b>18.160</b>	10.430	4.250	4.736
		GreenMarl	<b>4.289</b>	3.492	<b>5.504</b>	<b>2.444</b>	<b>2.218</b>	1.173	18.823	<b>9.602</b>	3.187	<b>3.573</b>
		StarPlat	6.530	11.060	7.650	3.300	4.422	1.777	22.250	11.740	5.580	12.83
	80	Galois	-	<b>7.269</b>	50.171	13.548	12.139	5.596	258.447	75.969	<b>7.435</b>	85.534
		Ligra	48.700	48.760	24.100	14.230	16.160	4.550	<b>73.230</b>	40.500	10.900	19.130
		GreenMarl	<b>18.189</b>	12.172	<b>21.094</b>	<b>9.353</b>	<b>8.578</b>	<b>4.107</b>	75.450	<b>39.704</b>	7.863	<b>14.18</b>
		StarPlat	33.240	44.270	32.150	13.050	16.550	6.810	86.400	47.290	18.030	50.260
	150	Galois	-	<b>14.598</b>	78.632	25.403	23.816	11.132	-	138.052	12.365	116.70
		Ligra	83.400	90.360	44.630	26.260	31.500	8.996	<b>124.66</b>	76.530	17.700	35.800
		GreenMarl	<b>32.990</b>	23.110	<b>38.612</b>	<b>17.709</b>	<b>17.231</b>	<b>8.184</b>	139.090	<b>74.506</b>	<b>11.069</b>	<b>26.18</b>
		StarPlat	61.625	84.430	57.664	24.520	32.290	13.553	160.865	80.111	32.136	94.540
PR	Galois	-	<b>0.510</b>	<b>0.647</b>	<b>0.371</b>	<b>0.474</b>	<b>0.156</b>	<b>0.607</b>	<b>0.224</b>	<b>0.324</b>	<b>0.44</b>	
	Ligra	25.600	162.660	5.050	3.930	3.623	0.836	2.050	0.822	5.880	0.942	
	GreenMarl	<b>0.585</b>	7.211	1.437	0.512	0.585	0.263	1.235	0.525	0.821	0.688	
	StarPlat	1.752	9.002	1.213	0.473	0.509	0.236	1.600	0.667	0.883	0.619	
SSSP	Galois	<b>0.522</b>	<b>0.132</b>	<b>0.404</b>	<b>0.203</b>	<b>0.205</b>	<b>0.099</b>	<b>19.387</b>	<b>6.798</b>	<b>0.133</b>	<b>0.480</b>	
	Ligra	10.7	0.148	5.136	1.846	3.89	1.683	283.000	9.043	2.74	10.40	
	GreenMarl	2.182	0.891	1.048	1.16	0.761	0.292	193.548	48.349	0.464	1.361	
	StarPlat	5.831	1.437	1.850	1.759	2.412	0.846	<b>294.303</b>	<b>53.395</b>	1.369	5.237	
TC	Galois	-	<b>56.432</b>	33.110	<b>46.168</b>	9.811	3.008	0.061	<b>0.020</b>	184.260	2.350	
	Ligra	2103.333	188.660	<b>22.800</b>	97.360	10.460	1.926	0.147	0.0698	<b>130.330</b>	1.706	
	GreenMarl	11611.029	4257.498	137.559	4564.568	29.426	12.705	0.065	0.021	5647.156	1.435	
	StarPlat	1414.323	59.925	23.420	111.430	<b>7.544</b>	<b>1.559</b>	<b>0.059</b>	0.024	158.760	<b>1.17</b>	

Table 3: StarPlat’s OpenMP code performance comparison against Galois, Ligra and Green-Marl (20 Threads). All times are in seconds. BC is run with the number of sources as {20, 80, 150}.

SSSP	TW	SW	OK	WK	LJ	PK	US	GR	RM	UR
	4.127	0.127	1.503	0.633	2.315	0.822	72.654	9.641	1.319	4.477

Table 4: StarPlat’s SSSP OpenMP code running times (seconds) with `static` scheduling


829 StarPlat-generated code) outperforms all the other benchmarks. This is due  
 830 to the in-place update of the PR values for vertices, which leads to faster  
 831 convergence. StarPlat, Green-Marl, and Ligra follow a similar processing of  
 832 updating the PR values using double buffering.

833 **Single-Source Shortest Paths.** We find that Galois SSSP is faster than  
 834 StarPlat’s code and other compared benchmarks. This is due to application-  
 835 specific prioritized scheduling in the Galois framework [? ]. For instance,  
 836 processing tasks in the ascending distance order reduces the total amount of  
 837 extra work done. GreenMarl and StarPlat have nearly similar implementa-  
 838 tions for SSSP calculation. Both follow a dense push configuration for vertex  
 839 processing which requires iterating over all the vertices to check if they are  
 840 active. This is typically costly for road networks which have a smaller fron-  
 841 tier active in each iteration. In addition, this additional cost accumulates  
 842 over several iterations due to the large diameter of road networks. However,  
 843 GreenMarl performs better than StarPlat for nearly all graphs. GreenMarl  
 844 uses spin-lock implementation with the back-off strategy to save on unnec-  
 845 essary CPU cycles. StarPlat also uses a lock-free atomic-based implementa-  
 846 tion but the unnecessary updates are more profound leading to performance  
 847 degradation due to false sharing. Ligra’s performance is not very competi-  
 848 tive with other benchmarks. Ligra switches between sparse and dense edge  
 849 processing based on the frontier size. But this direction optimization does  
 850 not lead to considerable performance improvement for the given graph suite  
 851 except for the road networks.


852 By default, StarPlat generates OpenMP code with `dynamic` scheduling.  
 853 This largely works well across various algorithms and graph types. However,  
 854 SSSP code seems to overall perform better with static scheduling as shown  
 855 in Table 4. The difference is pronounced for large diameter graphs US and  
 856 GR wherein the execution times reduce from over a minute to a few seconds.

857 **Triangle Counting.** Galois, StarPlat, and Green-Marl follow a node-  
 858 iterator pattern in TC. On the other hand, Ligra follows an edge-iterator  
 859 based version, which is supposed to work better for skewed degree graphs,  
 860 since the edge-based version has better load balance. We observe that for  
 861 different graphs, different frameworks outperform. Interestingly, performance  
 862 of the Green-Marl-generated code is significantly poorer.

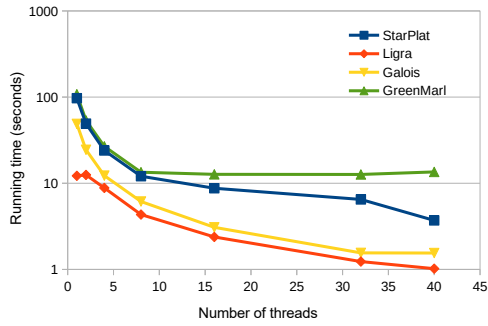





(a) UR



(b) GR



(c) PK



(d) LJ

Figure 17: Log plots of running times of various types of graphs with varying number of threads (note different y-axis scales).

Algo.		Framework	TW	SW	OK	WK	LJ	PK	US	GR	RM	UF
BC	20	Galois	2814	3502	1733	54776	2962	2331	>95 Hrs	>12 Hrs	1006	26
		StarPlat	<b>442</b>	<b>13</b>	<b>681</b>	<b>244</b>	<b>197</b>	<b>52</b>	<b>917</b>	<b>353</b>	<b>437</b>	<b>7</b>
	80	Galois	12375	14043	6999	>12 Hrs	11320	9035	>12 Hrs	>12 Hrs	2479	1024
		StarPlat	<b>2102</b>	<b>49</b>	<b>2705</b>	<b>978</b>	<b>764</b>	<b>194</b>	<b>3741</b>	<b>1489</b>	<b>975</b>	<b>29</b>
	150	Galois	NA	NA	NA	NA	NA	NA	NA	NA	NA	N
		StarPlat	<b>3816</b>	<b>94</b>	<b>4973</b>	<b>1829</b>	<b>1498</b>	<b>382</b>	<b>6782</b>	<b>2713</b>	<b>1442</b>	<b>55</b>
PR		Galois	398	761	314	376	551	479	257	296	244	20
		StarPlat	<b>57</b>	OOM	<b>23</b>	<b>11</b>	<b>8</b>	<b>2</b>	<b>2</b>	<b>0.3</b>	<b>44</b>	
SSSP		Galois	<b>33</b>	11	<b>30</b>	63	<b>27</b>	24	<b>1933</b>	82	<b>20</b>	4
		StarPlat	152	<b>0.4</b>	104	<b>35</b>	47	<b>15</b>	2302	<b>27</b>	70	<b>1</b>
TC		Galois	<b>71275</b>	<b>436</b>	<b>327</b>	<b>884</b>	<b>174</b>	<b>19</b>	7	9	<b>1238</b>	<b>1</b>
		StarPlat	>24 Hrs	>72 Hrs	64114	>24 Hrs	5741	2420	<b>4</b>	<b>0.6</b>	>24 Hrs	<b>3</b>

Table 5: StarPlat’s MPI code performance comparison against Galois (96 processes). All times are in seconds. BC is run with the number of sources as {20, 80, 150}. StarPlat goes out-of-memory on SW in PR due to double buffering.

## 863 5.2. MPI

864 The comparison of execution time of StarPlat’s MPI code against the  
865 distributed Galois framework for four algorithms is given in Table 5. The  
866 execution time does not include the graph reading and distribution time.  
867 OpenMPI 3.1.6 is used for execution and the algorithms are executed with  
868 96 number of processes.

869 **Betweenness Centrality.** From the table, we can see that, StarPlat’s BC  
870 code takes much lesser time for execution compared to the Galois framework  
871 for all the ten graphs. When 20 number of sources are used, the execution  
872 of Galois code for two graphs TW and GR is not completed even after 96  
873 hours and 12 hours respectively. Similarly, for 80 number of sources, the  
874 Galois code does not give result for three graphs WK, US and GR till 12  
875 hours of execution. The execution of Galois code for all the 10 graphs with  
876 150 number of sources is not tried as it is expected to take much more time  
877 than StarPlat’s BC code for all the ten graphs based on the execution time  
878 for 20 and 80 number of sources.

879 **PageRank.** StarPlat-generated PR code outperforms the Galois PR code  
880 for all the graphs except SW. StarPlat’s PR on SW goes out of memory due  
881 to double-buffering. We are currently investigating a way to fix it.

882 **Single-Source Shortest Paths.** From the comparison of execution time

Algo.		Framework	TW	SW	OK	WK	LJ	PK	US	GR	RM	UR
BC	1	LonestarGPU	-	-	-	-	-	-	-	-	-	-
	1	Gunrock	2.122	4.237	0.525	0.535	0.548	0.317	<b>2.811</b>	<b>1.750</b>	1.238	0.944
	1	StarPlat	<b>0.002</b>	<b>0.004</b>	<b>0.149</b>	<b>0.153</b>	<b>0.078</b>	<b>0.029</b>	17.656	6.359	<b>0.225</b>	<b>0.079</b>
	20	StarPlat	6.992	2.279	2.762	3.014	1.298	0.534	369.701	126.485	2.949	1.593
	80		28.179	9.332	11.331	12.050	4.886	1.907	1444.656	518.968	6.509	6.372
	150		55.548	MLE	21.241	27.271	9.609	3.794	2636.453	978.758	9.912	11.957
PR		LonestarGPU	-	<b>0.240</b>	0.363	<b>0.104</b>	<b>0.225</b>	<b>0.240</b>	<b>0.832</b>	<b>0.294</b>	<b>0.240</b>	<b>0.240</b>
		Gunrock	15.230	36.910	2.430	2.460	2.952	1.085	13.345	6.499	9.170	5.487
		StarPlat	<b>4.081</b>	7.112	<b>0.256</b>	1.780	1.300	0.257	3.420	0.679	0.891	0.257
SSSP		LonestarGPU	0.045	0.077	0.217	0.058	0.084	0.037	<b>0.162</b>	<b>0.091</b>	0.129	0.183
		Gunrock	2.272	4.057	0.616	0.556	0.562	0.311	1.283	1.140	1.034	0.915
		StarPlat	<b>0.001</b>	<b>0.002</b>	<b>0.078</b>	<b>0.044</b>	<b>0.027</b>	<b>0.012</b>	1.667	0.695	<b>0.120</b>	<b>0.028</b>
TC		LonestarGPU	-	31.990	2.998	2.771	<b>0.110</b>	<b>0.039</b>	11.874	5.695	<b>1.270</b>	0.499
		Gunrock	<b>67.718</b>	7.369	<b>0.843</b>	<b>0.997</b>	0.850	0.404	1.490	0.712	3.200	1.040
		StarPlat	10540.002	<b>1.410</b>	46.700	4.009	3.006	0.655	<b>0.001</b>	<b>0.001</b>	824.620	<b>0.034</b>

Table 6: StarPlat’s CUDA code performance comparison against LonestarGPU and Gunrock. All times are in seconds. LonestarGPU does not have BC implemented and fails to load the largest graph TW.

of SSSP code of StarPlat with Galois, we can observe that, Galois performs better for 5 graphs TW, OK, LJ, US and RM while StarPlat performs better for the remaining five graphs. The number of iterations taken by StarPlat is fewer than that by Galois. This is due to multiple levels of distance update within the local nodes of a process in a single outer loop iteration.

**Triangle Counting.** The TC code generated by StarPlat is less efficient as it iterates through neighbours of each node multiple times and check for neighbourhood among these nodes. Due to this, StarPlat’s TC code takes much more time than Galois. For road graphs US and GR, StarPlat’s TC takes lesser time than Galois code. For some graphs such as TW, SW, WK and RM StarPlat’s code did not give result even after 24 hours.

### 5.3. CUDA

Table 6 presents the absolute running times of the four algorithms on our ten graphs for the three frameworks: LonestarGPU, Gunrock, and StarPlat. The running times include CPU-GPU data transfer.

**Betweenness Centrality.** StarPlat-generated code outperforms Gunrock’s library-based code on eight out of ten graphs. On the two road

900 networks, Gunrock is considerably better than StarPlat. We use CUDA’s  
 901 `grid.synchronization` to our benefit, whereas Gunrock relies heavily on  
 902 this bulk-synchronous processing. Gunrock’s Dijkstra’s algorithm works very  
 903 well for road networks. On the other hand, on social and random graphs,  
 904 our implementation fares better. We also illustrate performance with mul-  
 905 tiple sources of different sizes (20, 80, and 150). Except for soc-sinaweibo,  
 906 StarPlat-generated code is on par with or better than the other frameworks.  
 907 Finally, unlike Gunrock and LonstarGPU, StarPlat has the provision to ex-  
 908 ecute BC from a set of sources.

909 **PageRank.** We observe that the three frameworks have consistent rela-  
 910 tive performance, with hand-crafted LonestarGPU codes outperforming the  
 911 other two and StarPlat outperforming Gunrock. StarPlat exploits the double  
 912 buffering approach to read the current PR values and generate those for the  
 913 next iteration. This separation reduces synchronization requirement during  
 914 the update, but necessitates a barrier across iterations. LonestarGPU uses  
 915 an in-place update of the PR values and converges faster.

916 **Single-Source Shortest Paths.** Gunrock uses Dijkstra’s algorithm for  
 917 computing the shortest paths using a two-level priority queue. We have coded  
 918 a variant of the Bellman-Ford algorithm in StarPlat. Hence, the comparison  
 919 may not be most appropriate. But we compare the two only from the appli-  
 920 cation perspective – computing the shortest paths from a source in the least  
 921 amount of time. LonestarGPU and StarPlat outperform Gunrock on all the  
 922 ten graphs. Between LonestarGPU and StarPlat, there is no clear winner.  
 923 They, in fact, have quite similar execution times.

924 **Triangle Counting.** Unlike other three algorithms, TC is not a propaga-  
 925 tion based algorithm. In addition, it is characterized by a doubly-nested loop  
 926 inside the kernel (see Figure 20). Another iteration is required for checking  
 927 edge (Line 7) which can be implemented linearly or using binary search if  
 928 the neighbors are sorted in the CSR representation. Due this variation in  
 929 the innermost loop, the time difference across various implementations can  
 930 be pronounced, which we observe across the three frameworks. Their perfor-  
 931 mances are mixed across the ten graphs, and no one emerges as an overall  
 932 winner.

## 933 6. Related Work

934 We divide the relevant related work based on the target backend in the  
935 next three subsections, and discuss generic frameworks in Section 6.4. .

### 936 6.1. *OpenMP*

937 Ligra [?] is a lightweight graph processing framework that is specific for  
938 shared-memory parallel/multi-core machines, which makes graph traversal  
939 algorithms easy to write. The framework has two very simple routines, one  
940 for mapping over edges and one for mapping over vertices. The routines  
941 can be applied to any subset of the vertices, which makes the framework  
942 promising for many graph traversal algorithms operating on subsets of the  
943 vertices. In abstraction, a vertex subset is maintained as a set of integer  
944 labels for the included vertices, and the routine for mapping over vertices  
945 applies the user-supplied function to each integer.

946 A lightweight infrastructure for graph analytics, Galois [?] supports a  
947 more general programming model which is more expressive than restrictive.  
948 The infrastructure supports autonomous scheduling of fine-grained tasks with  
949 application-specific priorities and a library for scalable data structures. Ex-  
950 isting DSLs can be implemented on top of the Galois system to achieve high  
951 throughput. One can think of a Galois backend for StarPlat.

952 Green-Marl [?] is a DSL for shared memory graph processing. Programs  
953 depending on BFS/DFS traversals can be written concisely with the corre-  
954 sponding constructs in the language. However, for other graph algorithms,  
955 the user has to define the iteration over vertices or edges explicitly. So, the al-  
956 gorithmic description can be written intuitively using the constructs provided  
957 by the language, while exposing the inherent data-level parallelism. StarPlat  
958 borrows a few ideas from the language, and also quantitatively compares its  
959 performance against Green-Marl, Ligra, and Galois.

960 GraphIt [?] is a DSL for graph analysis that achieves high performance  
961 by enabling programmers to easily find the best combination of optimiza-  
962 tions for their specific algorithm and input graph. It essentially separates  
963 computation from scheduling. The algorithms are specified using an algo-  
964 rithm language, and performance optimizations are specified using a separate  
965 scheduling language.

### 966 6.2. *MPI*

967 D-Galois, a distributed graph processing system, was formed by inter-  
968 facing Gluon, a communication optimization technique with Galois shared

memory graph processing system [? ]. We compare against Gluon in our experiments.

Pregel and Giraph graph processing frameworks were inspired by the BSP programming model [? ? ]. Pregel decomposes the graph processing computation as a sequence of iterations in which the vertices receive data sent in the previous iteration, changes its state, and sends messages to its out neighbors which will be received in the next iteration. GPS provided additional features such as enabling global computation, dynamic graph repartitioning, and repartitioning of large adjacency lists to the Pregel framework [? ].

A compiler automatically generates Pregel code from a subset of programs called Pregel canonical written using Green-Marl DSL [? ]. It also applies certain transformations such as edge flipping, dissecting nested loops, translating random access, transforming BFS and DFS traversals to convert non-Pregel canonical pattern into Pregel canonical. Another compiler named DisGCo can translate all the Green-Marl programs to equivalent MPI RMA programs [? ]. It handled the challenges such as syntactic differences, differences in memory view, intermixed serial code with parallel code, and graph representation while translating a Green-Marl program to MPI RMA code.

The specification and implementation of morph graph algorithms and non-vertex centric graph algorithms for heterogeneous distributed environments was provided by the DH-Falcon DSL and its compiler [? ]. A hybrid approach that uses both dense and sparse mode processing by utilizing pull- and push-based computation approaches was provided by a distributed graph analytics scheme Gemini [? ]. Additionally, it used optimized graph representation and partitioning methods for intra-node as well as inter-node load balancing.

GraphLab provides a fault-tolerant distributed graph processing abstraction that can perform a dynamic, asynchronous graph processing for machine learning algorithms [? ]. Another parallel distributed graph processing abstraction for processing power law graphs with a caching mechanism and a distributed graph placement mechanism was introduced by PowerGraph [? ]. GraphChi supports asynchronous processing of dynamic graphs using a sliding approach on the disk-based systems [? ].

### 6.3. CUDA

Gunrock [? ] is a graph library which uses data-centric abstractions to perform operations on edge and vertex frontier. All Gunrock operations are bulk-synchronous, and they affect the frontier by computing on values within

it or by computing a new one, using the following three functions: *filter*, *compute*, and *advance*. Gunrock library constructs efficient implementations of frontier operations with coalesced accesses and minimal thread divergence.

LonestarGPU [?] is a collection of graph analytic CUDA programs. It employs multiple techniques related to computation, memory, and synchronization to improve performance of the underlying graph algorithms. We quantitatively compare StarPlat against Gunrock and LonestarGPU.

Medusa [?] is a software framework which eases the work of GPU computation tasks. Similar to Gunrock, it provides APIs to build upon, to construct various graph algorithms. Medusa exploits the BSP model, and proposes a new model EVM (Edge Message Vertex), wherein the local computations are performed on the vertices and the computation progresses by passing messages across edges.

CuSha [?] is a graph processing framework that uses two graph representations: G-Shards and Concatenated Windows(CW). G-Shards makes use of a recently developed idea for non-GPU systems that divides a graph into ordered sets of edges known as shards. In order to increase GPU utilisation for processing sparse graphs, CW is a novel format that improves the use of shards. CuSha makes the most of the GPU’s processing capability by processing several shards in parallel on the streaming multiprocessors. CuSha’s architecture for parallel processing of large graphs allows the user to create the vertex-centric computation and plug it in, making programming easier. CuSha significantly outperforms the state-of-the-art virtual warp-centric approach in terms of speedup.

MapGraph [?] is a parallel graph programming framework which provides a high level abstraction which helps in writing efficient graph programs. It uses SOA(Structure Of Arrays) to ensure coalesced memory access. It uses the dynamic scheduling strategy using GAS(Gather-Apply-Scatter) abstraction. Dynamic scheduling improves the memory performance and dispense the workload to the threads in accordance with degree of vertices.

#### 6.4. Generic Frameworks

Kakkos 3 [?] is a programming model in C++ for writing performance portable applications targeting all major HPC platforms (CUDA, HIP, SYCL, HPX, OpenMP and C++ threads as backend). It provides abstractions for both parallel executions of code and data management. However, Kokkos kernels are used mainly on matrix-based operations or algorithms, and it targets coloring algorithms only recently, and many other complex

1043 algorithms are not supported (as of date). Philosophy-wise, Kokkos and  
1044 StarPlat have similarity, while the former is generic and is not a DSL.

1045 Graphit [?] adds support for code generation on GPUs along with CPUs  
1046 in their recent version. It expands the optimization space of GPU graph  
1047 processing frameworks with a novel GPU scheduling language, and a compiler  
1048 which enables various optimizations such as combining load balancing, edge  
1049 traversal direction, active vertex set creation, active vertex set processing  
1050 ordering, and kernel fusion.

1051 Julia [?] is a high-level, dynamic programming language designed for  
1052 high performance. It is well suited for numerical analysis and computational  
1053 science. It supports concurrent, parallel, and distributed computing (with  
1054 and without MPI). Circle [?] is a new C++ 20 compiler. It extends C++  
1055 by adding many novel language features. Circle is a heterogeneous compiler  
1056 that supports GPU compute and shader programming as well. With Circle,  
1057 one can write a device-portable GPGPU code with Vulkan compute (using  
1058 real pointers into physical storage buffers), a powerful feature only imple-  
1059 mented by the Circle compiler. Carbon [?] is an experimental successor  
1060 language of C++ in development. One of its main goals is to support mod-  
1061 ern OS platforms, hardware architectures, environments, and fast & scalable  
1062 development (similar to Python and Rust). Odin [?] is again a general-  
1063 purpose programming language with distinct typing built especially for high-  
1064 performance, modern systems and data-oriented programming. It is under  
1065 development and supports structure of arrays (SoA) data types and array  
1066 programming. Chapel [?] is a modern parallel programming language  
1067 aimed at portability and scalability. The primary goal of Chapel is to sup-  
1068 port general parallel programming and also make parallel programming at a  
1069 scale far more productive.

1070 The recent sprouts of the above languages motivate the need for a sim-  
1071 ple language that is portable across different parallel programming platforms  
1072 and architectures without compromising the productivity of the users, irre-  
1073 spective of the domain they work for.

1074 In a similar spirit, Intel [?] and Nvidia [?] are venturing into hetero-  
1075 geneous programs which converts simple C++ program to parallel programs  
1076 running on various architectures (multi-cores, GPUs, FPGAs and even Arm  
1077 Servers). However, such programs are parallelized only if the program uses  
1078 STL algorithms. Their compiler replaces the standard C++'s STL algo-  
1079 rithms with the corresponding parallel versions based on the target. Intel  
1080 OneAPI (which includes DPC++ and TBB) uses C++ STL, Parallel STL



(PSTL), Boost Compute, and SYCL for parallelization. Nvidia’s `nvc++` uses a different linker for the target processors using command-line arguments. `nvc++` supports ISO C++17, and targets GPU and multicore CPU programming with C++17 parallel algorithms, OpenACC, and OpenMP.

We observe that Gunrock’s Essentials-cpp and Essentials [?] are along similar lines and use modern C++, `std` parallelism constructs and lambdas built with their bulk-synchronous-asynchronous, data-centric abstraction model. They currently implement BFS and SSSP.

## 7. Conclusion

We presented StarPlat, a DSL for implementing graph analytic algorithms. The language model provides various constructs for expressing graph problems at a high level without the complexities of implementing the commonly occurring patterns from scratch. This also facilitates providing optimized solutions both in terms of memory and time for these patterns. Currently, we support development across multicore, distributed, and manycore backends. We discussed the translation of StarPlat’s constructs for various backends. With a range of large graphs and four different algorithms, we experimentally validated that the performance of the generated code is comparable to the hand-tuned and generated implementations of existing frameworks and libraries.

While there are multiple software solutions available in the form of programming languages and libraries for the development of graph algorithms, an important next step is to build solutions that are versatile to support multi-platform development. StarPlat is an effort in this direction. The future direction of StarPlat is to improve the portability of the solution to support OpenACC and OpenCL backends, and to enable automatic code generation for heterogeneous architectures.

## Acknowledgments

We gratefully acknowledge the use of the computing resources at HPCE, IIT Madras. This work is supported by India’s National Supercomputing Mission grant CS1920/1123/MEIT/008606.

## 1112 Appendix A. Graph Algorithms in StarPlat

1113 StarPlat specifications for Betweenness Centrality, PageRank, and Tri-  
 1114 angle Counting are shown below in Figures 18, 19, and 20 respectively. In  
 1115 addition, Figure 21 shows pull-based SSSP.

```

1116 1 function computeBC(Graph g, propNode<float> BC,
1117   SetN<g> sourceSet) {
1118 2   g.attachNodeProperty(BC = 0);
1119 3   for (src in sourceSet) {
1120 4     propNode<double> sigma;
1121 5     propNode<float> delta;
1122 6     g.attachNodeProperty(delta = 0);
1123 7     g.attachNodeProperty(sigma = 0);
1124 8     src.sigma = 1;
1125 9
1126 10    iterateInBFS(v in g.nodes() from src) {
1127 11      for (w in g.neighbors(v)) {
1128 12        w.sigma = w.sigma + v.sigma;
1129 13      } }
1130 14
1131 15    iterateInReverse(v != src) {
1132 16      for (w in g.neighbors(v)) {
1133 17        v.delta = v.delta + (v.sigma /
1134 18          w.sigma) * (1 + w.delta);
1135 19      }
1136 20      v.BC = v.BC + v.delta;
1137 21    } } }

```

Fig 18: BC computation in StarPlat

```

1138 1 function computePR(Graph g, float beta, float
1139   delta, int maxIter, propNode<float> pageRank)
1140 {
1141 2   float num_nodes = g.num_nodes();
1142 3   propNode<float> pageRank_nxt;
1143 4   g.attachNodeProperty(pageRank = 1 /
1144   num_nodes);
1145 5   int iterCount = 0;
1146 6   float diff;

```

```

1147     7
1148     8     do {
1149     9         diff = 0.0;
1150    10
1151    11         forall(v in g.nodes()) {
1152    12             float sum = 0.0;
1153    13
1154    14             for (nbr in g.nodes_to(v)) {
1155    15                 sum = sum + nbr.pageRank /
1156    16                     g.count_outNbrs(nbr);
1157    17             }
1158    18
1159    19             float val = (1 - delta) / num_nodes
1160    20                 + delta * sum;
1161    21             diff += val - v.pageRank;
1162    22             v.pageRank_nxt = val;
1163    23         }
1164    24
1165    25         pageRank = pageRank_nxt;
1166    26         iterCount++;
1167    27     } while ((diff > beta) && (iterCount <
1168    28         maxIter));
1169
1170    29 }

```

Fig 19: PR computation in StarPlat

```

1171     1 function computeTC(Graph g) {
1172     2     int triangle_count = 0;
1173     3
1174     4     forall(v in g.nodes()) {
1175     5         forall(u in g.neighbors(v).filter(u <
1176     6             v)) {
1177     7             forall(w in g.neighbors(v).filter(w
1178     8                 > v)) {
1179     9                 if (g.is_an_edge(u, w)) {
1180     10                     triangle_count += 1;
1181     11                 }
1182     12             }
1183     13         }
1184     14     }
1185     15 }

```

Fig 20: TC computation in StarPlat

```

1182 1 function computePullSSSP (Graph g, node src) {
1183 2     propNode<int> dist;
1184 3     propNode<bool> modified;
1185 4     g.attachNodeProperty(dist = INF, modified =
1186         False);
1187 5     src.modified = True;
1188 6     src.dist=0;
1189 7     bool finished = False;
1190 8     fixedPoint until (finished: !modified) {
1191 9         forall (v in g.nodes()) {
1192 10             forall (nbr in
1193                 g.nodes_to(v).filter(modified ==
1194                     True) {
1195 11                 edge e = g.getEdge(v, nbr);
1196 12                 <v.dist, v.modified> =
1197                     <Min(v.dist, nbr.dist +
1198                         e.weight), True>;
1199 13 }     }     }     }

```

Fig 21: SSSP-Pull computation in StarPlat